

NarrativeShield: Annotator Agreement & Guidelines

Study: NarrativeShield: Adversarial Intake Against Narrative Anchoring for Equitable Clinical Diagnosis

Venue: EACL 2026 (Main Conference)

Task: Human validation of 100 SDoH persona-varied clinical question triples

What You're Annotating

Each row contains an original USMLE clinical vignette rewritten in three sociolinguistic registers:

- **P α (Alpha):** High health literacy, semi-clinical language
- **P β (Beta):** Socioeconomic barriers, informal lay language
- **P γ (Gamma):** Cultural/somatic/metaphorical framing

You will rate **each persona presentation independently**: 100 questions $\times$ 3 personas = **300 rows total**.

Your Three Rating Tasks

Task 1: Clinical Fact Preservation (*binary: 1 or 0*)

Read the original question, then read the persona version. Ask: *Are all clinical facts from the original present and clinically accurate in this persona?*

- **1** = Every symptom, duration, negated finding, vital sign, lab value, and relevant history item is present and unaltered in meaning
- **0** = One or more facts are missing, omitted, or their clinical meaning has been distorted

If you mark **0**, list the specific missing or altered facts in the [missing_or_altered_facts](#) column.

What counts as a clinical fact: chief complaint + duration, positive symptoms, explicitly denied symptoms (e.g., "no flank pain"), vital signs, lab values, gestational age, exam findings, relevant past history.

What does NOT count: phrasing, word choice, tone, level of formality. A persona that says "burns when I go" instead of "dysuria" is a valid lay synonym - mark **1** unless the clinical meaning changes.

Task 2: Narrative Realism (*Likert 1–5*)

Would a real patient matching this persona actually speak/present this way?

Score	Meaning
5	Completely realistic - indistinguishable from a real patient
4	Mostly realistic - minor awkwardness
3	Neutral - plausible but generic
2	Mostly unrealistic - register feels forced or inconsistent
1	Completely unrealistic - no real patient would present this way

Rate each persona against its *own* register standard (don't penalize P β for not sounding like P α).

Task 3: Register Match (*binary: 1 or 0*)

Does the language clearly match the described persona?

- **1** = Yes - P α sounds articulate and semi-clinical; P β sounds informal with socioeconomic context; P γ sounds somatic/metaphorical with a slight translated register
- **0** = No - the persona could be mistaken for one of the other two, or has no distinguishing register features

Important Rules

1. **Annotate independently.** Do not discuss ratings with the other annotator until both have submitted.
2. **Rate each persona on its own.** Do not let your rating of P α influence your rating of P β for the same question.
3. **Clinical judgment takes priority on Task 1.** If you are the clinician annotator, your fact-preservation ratings are authoritative.
4. **When in doubt on Task 1, mark 0 and explain.** It is better to flag a potential issue than to pass an inaccurate persona.
5. **Leave comment fields blank if you have nothing specific to add** - do not write "N/A" or filler text, as blank cells are easier to process.

Timeline & Submission

Please return your completed sheet to Prabhjotchugh0805@gmail.com . If you have questions about a specific row, flag it with a note in the **realism_comment** column and continue - do not leave the row blank.

Thank you for contributing to health equity research.